

Couchbase DB

Aktuelle DB Architekturen und Technologien

Vorlesung des Studienganges Informatik
an der Dualen Hochschule Baden-Württemberg Stuttgart

von

Robert Orbach, Oliver Mohr, Oliver Braun

28.3.2025

Inhaltsverzeichnis

1	Type of Database & History	1
1.1	Key Characteristics and Features	1
1.1.1	Data Model and Storage	1
1.1.2	Query and Indexing	1
1.1.3	Distribution and Scalability	1
1.1.4	Security and Enterprise Features	1
1.2	Brief History	1
1.3	Use Case Examples	2
2	Which aspects of relational DBs are improved?	2
2.1	Schema Flexibility	2
2.2	Scalability Improvements	2
2.3	Advanced Query Capabilities	3
2.4	Relationship to Codd's Rules	3
3	Deployment and Installation	4
3.1	Cloud Deployment	4
3.1.1	Using Couchbase Capella	4
3.2	On-premise Installation	4
3.2.1	Docker Installation	4
3.3	Choosing the Right Deployment Method	5
4	APIs and SDKs	5
4.1	Couchbase SDKs	5
4.2	Couchbase REST API	5
4.3	Example Code - Python	5
5	Couchbase Architecture	7
5.1	Data Storage and Distribution	7
5.2	Memory-First Architecture	7
5.3	Querying with N1QL (SQL++)	7
5.4	Cross Data Center Replication (XDCR)	7
5.5	Specialized Services in Couchbase	7
5.6	Multi-Dimensional Scaling (MDS)	8
5.7	CAP-Theorem	8
6	Advantages and Disadvantages	8
6.1	Technical Advantages	8
6.2	Technical Limitations	9
6.3	Technical Comparison	9

Literaturverzeichnis	10
---------------------------------------	-----------

1 Type of Database & History

Couchbase is a distributed NoSQL database that combines document database and key-value store capabilities. Unlike traditional single-model NoSQL solutions, Couchbase offers a hybrid approach that makes it versatile for different use cases.

1.1 Key Characteristics and Features

Couchbase distinguishes itself through several core capabilities that combine to create a versatile database platform:

1.1.1 Data Model and Storage

- **JSON Document Storage:** Schema-flexible documents supporting nested attributes, arrays, and objects (up to 20MB per document)
- **Key-Value Operations:** Sub-millisecond CRUD operations with direct key access, optimistic concurrency (CAS), and TTL expiration
- **Memory-First Architecture:** Built-in caching layer with configurable memory quotas and background disk persistence

1.1.2 Query and Indexing

- **SQL++/N1QL Query Language:** SQL-like syntax for JSON with support for JOINS, nested attributes, arrays, aggregations, and subqueries
- **Comprehensive Indexing:** Primary, secondary, composite, partial, and array indexes to optimize query performance

1.1.3 Distribution and Scalability

- **Distributed Architecture:** Automatic sharding, configurable replication, auto-failover, and online rebalancing
- **Multi-Service Design:** Independent scaling of data, query, index, search, analytics, and eventing services
- **Cross-Datcenter Replication:** Built-in XDCR for geographic distribution

1.1.4 Security and Enterprise Features

- **Access Control:** Role-Based Access Control (RBAC), LDAP integration, and operation auditing
- **Encryption:** TLS for data in transit and field-level encryption for sensitive data

1.2 Brief History

Couchbase was formed in 2010 through the merger of Membase (a key-value store with Memcached compatibility) and CouchDB technology (document database capabilities). Key milestones include:

- 2011: Release of Couchbase Server 1.0
- 2015: Introduction of N1QL, a SQL-like query language for JSON
- 2017: Launch of Couchbase Mobile for edge computing
- 2021: Couchbase goes public with IPO

The introduction of N1QL in 2015 was particularly significant, as it bridged the gap between NoSQL flexibility and SQL familiarity.

1.3 Use Case Examples

Couchbase excels in several common application scenarios:

User Profile Management: Document model for dynamic attributes, fast key-value access, and session expiration.

Product Catalogs: Flexible schema for product attributes, N1QL for advanced queries, and full-text search.

Gaming Applications: Low-latency key-value operations, document model for player/game data, and horizontal scaling.

Internet of Things: Mobile sync for edge devices, time-series support for sensor data, and schema flexibility for device diversity.

2 Which aspects of relational DBs are improved?

Couchbase addresses three fundamental limitations of relational databases: schema rigidity, scaling challenges, and query limitations for complex data structures.

2.1 Schema Flexibility

Relational limitation: Relational databases require predefined schemas where all data must conform to a fixed structure. Schema changes typically involve ALTER TABLE operations that often cause downtime, require data migrations, and necessitate application code updates.

Other NoSQL limitations:

- **MongoDB:** Requires field-level indexes specified in advance, with indexing limitations on deeply nested structures
- **Cassandra:** Uses table-like structures requiring column family definitions upfront
- **DynamoDB:** Limits secondary indexes and requires careful attribute planning

Couchbase advantage: Couchbase's document model delivers superior flexibility through:

- Schema-free JSON documents with unlimited nesting depth
- No predefined structure requirements at collection level
- Full indexing capabilities on any field or nested path
- Support for heterogeneous documents within the same collection
- Built-in schema validation when governance is needed

2.2 Scalability Improvements

Relational limitation: Traditional relational databases were designed for vertical scaling (bigger servers) rather than horizontal scaling. As data grows, this approach eventually hits hardware limits. Manual sharding requires complex application logic, and joins become problematic across distributed data.

Couchbase advantage: Couchbase provides a truly distributed architecture that surpasses both relational and other NoSQL solutions:

- **Linear Performance:** Near-linear throughput increases with additional nodes (95% efficiency vs. theoretical maximum)
- **In-Memory Optimization:** Memory-first architecture with configurable RAM-to-disk ratios
- **Intelligent Rebalancing:** Auto-weighted data redistribution during cluster changes
- **Multi-Region:** Advanced cross-datacenter replication (XDCR) with filtering and compression

2.3 Advanced Query Capabilities

Relational limitation: SQL struggles with hierarchical data structures common in modern applications. Complex nested structures must be split across multiple tables, requiring joins that become performance bottlenecks at scale. This creates a mismatch with object-oriented application code.

Couchbase solution: SQL++/N1QL (SQL for JSON) combines SQL familiarity with direct operations on nested JSON structures:

```
SELECT product.name, product.specs.cpu
FROM products AS product
WHERE "black" IN product.colors;
```

This query directly accesses nested fields and array elements without complex joins or subqueries.

2.4 Relationship to Codd's Rules

Couchbase strategically diverges from Codd's relational model while preserving key benefits and adding NoSQL advantages:

Principle	Relational Approach	MongoDB Approach	Couchbase Approach
Data Structure	Fixed tables and columns	BSON documents	JSON documents with optional schemas
Query Language	SQL	Proprietary	SQL++ (N1QL)
Transactions	ACID by default	ACID with limitations	Configurable ACID with versatile consistency
Relationships	Foreign keys and constraints	Manual reference or embedding	Flexible references with JOIN support
Schema Management	Strict schema enforcement	Schema-optional	Schema-optional with validation
Distribution Model	Complex federation	Replica sets with sharding	Unified cluster with multi-dimensional scaling

Couchbase fundamentally reimagines database architecture by combining the query power and familiarity of relational systems, the flexibility of document databases, the performance of key-value stores, and the scalability of distributed systems—creating a truly versatile data platform that eliminates the traditional tradeoffs between different database paradigms.

3 Deployment and Installation

Couchbase Server can be deployed in two different ways, depending on the user's needs and infrastructure preferences.

3.1 Cloud Deployment

The simplest way to use Couchbase Server is through Couchbase Capella, a fully managed cloud-based service. This eliminates the need for an on-premise installation and ongoing maintenance, making it an excellent choice for users who prefer a hassle-free, scalable solution.

3.1.1 Using Couchbase Capella

To get started with Couchbase Capella, these steps can be followed:

1. Sign up for a Couchbase Capella account by visiting the official Couchbase Capella website and following the registration instructions. (Couchbase 2025a)
2. Once the account is created, set up a new Couchbase Server cluster by following the provided step-by-step guide.
3. After the cluster is successfully deployed, a sample-bucket is already added, and a connection string and credentials can be created to access the cluster.

(Rottach 2021)

3.2 On-premise Installation

For users who prefer more control over their deployment, installing Couchbase Server (on-prem) is a viable option. The recommended approach for this is using Docker, which simplifies the setup process while offering flexibility.

3.2.1 Docker Installation

To install Couchbase Server with Docker, these steps can be followed:

1. (If not already done:) Install Docker on your machine by following the official installation guide available on the Docker website.
2. Pull and run the Couchbase Server container with the following command:

```
docker run -d --name db -p 8091-8097:8091-8097 -p 9123:9123 -p 11207:11207 -p 11210:11210 -p 11280:11280 -p 18091-18097:18091-18097 couchbase
```
3. Once the container is running, access the Couchbase Web Console by opening a web browser and navigating to <http://localhost:8091>
4. Follow the on-screen instructions to set up the Couchbase Server cluster and configure the necessary settings.

(Docker 2025)

3.3 Choosing the Right Deployment Method

Selecting between Couchbase Capella and an on-premise installation depends on various factors. Thus it is essential to consider the advantages and disadvantages of each deployment method before making a decision.

Couchbase Capella, as a cloud deployment, offers a fully managed service with automated updates and backups, ensuring high availability and scalability without the need for infrastructure maintenance. However, it comes with ongoing cloud usage costs and offers less flexibility in configuration and customization.

In contrast, an on-premise installation using Docker provides full control over configuration and security settings while avoiding cloud-related costs. Yet, it requires manual maintenance, updates, and backups, is limited by local machine resources, and involves a more complex initial setup compared to Capella.

Ultimately, Capella is ideal for businesses seeking a hassle-free, scalable solution, while an on-prem installation is better suited for developers who need full control over their environment.

4 APIs and SDKs

Couchbase provides a variety of APIs and SDKs to interact with the database, making it easy to integrate Couchbase into your applications. These APIs and SDKs are available in multiple programming languages, allowing developers to work with Couchbase using their preferred language.

4.1 Couchbase SDKs

Couchbase offers official SDKs for popular programming languages, including Java, .NET, Node.js, Python, Go, and others. These SDKs provide a high-level interface to interact with Couchbase, simplifying the development process and enabling developers to focus on building their applications. With the SDKs developers can perform all operations supported by Couchbase. (Couchbase 2025b)

4.2 Couchbase REST API

In addition to the SDKs, Couchbase also provides a REST API that allows developers to interact with the database over HTTP. This API is useful for scenarios where a native SDK is not available or when integrating with other systems that support RESTful communication.

4.3 Example Code - Python

The Python SDK provides a high-level interface to interact with Couchbase, making it easy to perform CRUD operations, query data, and manage the cluster. To install the couchbase module run:

```
pip install couchbase (brett19 u. a. 2025)
```

Let's look at a code snippet that demonstrates basic CRUD operations:

Skript 4.1: Python example

```
# Sample airline document
sample_airline = {
    "type": "airline",
    "id": 8091,
    "callsign": "CBS",
    "iata": None,
    "icao": None,
    "name": "Couchbase Airways",
}
key = "airline_8091"

# Create a document with specific ID
result = collection.insert(key, sample_airline)

# Retrieve a document by ID
result = collection.get(key)

# Update a document by ID
sample_airline["name"] = "Couchbase Airways!!"
result = collection.replace(key, sample_airline)

# Delete a document by ID
result = collection.remove(key)
```

We can see that interacting with Couchbase using the couchbase-python-module is straightforward and requires minimal code to perform common database operations.

5 Couchbase Architecture

Couchbase runs as a cluster of multiple server nodes, with each node able to take on different roles, such as data storage, indexing, or query processing. This distributed approach ensures a balanced load and makes the system more stable, as failures of individual nodes do not cause major issues.

5.1 Data Storage and Distribution

In Couchbase, data is stored in buckets, which serve as logical containers. To improve data management, these buckets are divided into smaller units called virtual buckets (vBuckets). This helps distribute data evenly across cluster nodes. Additionally, vBuckets allow dynamic data redistribution when nodes are added or removed, improving scalability and performance.

To ensure high availability and fault tolerance, Couchbase uses automatic data replication. Each vBucket has a primary copy and one or more replica copies stored on different nodes. If a node fails, a replica can automatically become the primary vBucket, keeping the system running without interruption.

5.2 Memory-First Architecture

Couchbase follows a Memory-First approach, where frequently used data and indexes are stored in RAM for quick access. Write operations happen in memory first and are later saved to disk or replicated to other nodes. This results in very low latency and significantly improves performance.

5.3 Querying with N1QL (SQL++)

N1QL, also called SQL++ is a query language that combines SQL-like syntax with the flexibility of JSON documents. This allows developers to run complex queries on NoSQL data while using familiar SQL structures. Couchbase supports advanced indexing to improve query performance and provides ACID transactions for data consistency.

5.4 Cross Data Center Replication (XDCR)

XDCR (Cross Data Center Replication) allows asynchronous data replication between different data centers or cloud environments. It improves global availability, supports disaster recovery, and keeps data synchronized across distributed applications. XDCR can work in one-way or two-way replication mode, giving flexibility for different needs.

To manage conflicts in two-way replication, XDCR offers conflict resolution strategies, such as timestamp-based resolution, where the latest update is kept, or custom conflict resolution, where developers define rules based on business logic.

5.5 Specialized Services in Couchbase

Couchbase Server provides several specialized services that can be deployed and managed separately. These services optimize data processing, improve scalability, and increase flexibility:

1. Data Service: Stores data in Key-Value format for fast read and write operations.
2. Query Service: Runs N1QL queries and optimizes them using indexes.

3. Index Service: Creates and manages indexes to speed up queries.
4. Search Service: Supports full-text search, including natural language processing and fuzzy matching, so that a search for "beautiesäalso finds "beautyänd "beautiful."
5. Analytics Service: Handles complex, long-running queries on large datasets without slowing down normal operations.
6. Eventing Service: Runs real-time server-side functions triggered by document changes or scheduled events.
7. Backup Service: Provides incremental and full backups, allowing data recovery and merging of previous backups.

5.6 Multi-Dimensional Scaling (MDS)

Couchbase uses Multi-Dimensional Scaling (MDS) to let each service scale independently on different nodes. Unlike traditional database systems, where every server does everything, MDS allows resources to be allocated based on specific needs. For example:

1. The Query Service can have more CPU and memory without affecting storage.
2. The Index Service can run on dedicated nodes for better performance.
3. The Data Service can scale separately to handle heavy read/write operations.

This separation improves performance, increases fault tolerance, and helps businesses adjust to changing workloads.

5.7 CAP-Theorem

Depending on the configuration Couchbase can either be considered a CP or an AP system. By default it leans more towards prioritizing high availability and performance making it a CP System. However, if it is configured to enforce stronger consistency mechanisms, for example ACID transactions, which are tunable regarding their durability and isolation levels. Morris 2023

6 Advantages and Disadvantages

6.1 Technical Advantages

Performance Metrics:

- Sub-millisecond key-value operations (<1ms @ 99th percentile)
- Memory-first architecture with configurable eviction policies
- B-tree based global secondary indexes for query optimization
- Write-optimized storage engine with append-only commits
- 30-40% storage reduction via snappy compression

Technical Capabilities:

- Multi-model support: K-V, document, spatial, full-text within single platform
- ANSI JOIN and NEST operations in N1QL with pushdown optimization
- Cross Datacenter Replication (XDCR) with filtering and compression
- SSLv3/TLS 1.2+ encryption with FIPS 140-2 compliance

- SDK support for Java, .NET, Node.js, Python, Go with reactive extensions

6.2 Technical Limitations

Performance Constraints:

- Document size limit: 20MB (default)
- Memory overhead: 56 bytes metadata per document
- Transaction latency: increased by 15-30% for multi-document ACID
- Query performance degrades with >10 JOINS in single statement
- Minimum 4GB RAM recommended per node for production

Architectural Considerations:

- Default consistency: eventual for queries, strong for K-V operations
- CAP theorem positioning: CP for document operations, AP for cross-cluster
- Minimum 3 nodes recommended for high availability
- Scaling requires rebalance operations (minimal but measurable impact)
- Analytics segregation required for OLAP without OLTP impact

6.3 Technical Comparison

Feature	Couchbase	MongoDB	Cassandra
Query Latency	1-5ms K-V, 5-20ms N1QL	2-10ms K-V, 10-50ms query	1-3ms K-V, 5-15ms CQL
Scaling Model	Shared-nothing, auto-sharding	Replica sets with sharding	Masterless ring
Consistency	Tunable (BASE to ACID)	Tunable (Read preferences)	Quorum-based
Indexes	Memory-optimized, composite	B-tree, compound	LSM-tree, materialized views
Transactions	ACID within/across docs	ACID within/across docs	Lightweight transactions

Ideal Workloads: High-throughput OLTP with sub-ms requirements; mixed document/K-V operations; distributed multi-model applications requiring SQL compatibility.

Less Suitable: Graph-centric applications (high relationship traversal); pure analytics workloads; single-server deployments with <8GB RAM.

Literaturverzeichnis

- brett19 u. a. (28. Jan. 2025). Couchbase PyPI. URL: <https://pypi.org/project/couchbase/> (besucht am 06.03.2025).
- Couchbase, Inc. (2025a). Couchbase Capella Sign Up. URL: <https://cloud.couchbase.com/sign-up> (besucht am 04.03.2025).
- (2025b). Couchbase SDKs. URL: <https://www.couchbase.com/developers/sdks/> (besucht am 06.03.2025).
- Docker, Inc. (2025). Dockerhub Couchbase. URL: https://hub.docker.com/_/couchbase (besucht am 06.03.2025).
- Morris, Jeff (2023). Data Consistency Models: Couchbase vs. CockroachDB. Blog-Post. URL: <https://www.couchbase.com/blog/consistency-models-couchbase-vs-cockroachdb/>.
- Rottach, Tim (18. Nov. 2021). Getting Started with Couchbase Capella in a Few Easy Steps. URL: <https://www.couchbase.com/blog/getting-started-with-couchbase-capella-in-a-few-easy-steps/> (besucht am 04.03.2025).